

# Real Time Automatic Modulation to Vitis

[Motivation](#)

[Testing With Python or HLS](#)

[Initial Implementation with Vitis HLS](#)

[Optimizing the design in Vitis HLS](#)

[Running In Vitis](#)

[Software and Hardware emulation outputs](#)

[Final Report files](#)

[Tradeoffs between Vitis and Vivado](#)

[Pros of the Vitis Unified Software Platform](#)

[Cons of the Vitis Unified Software Platform](#)

[Future Work \(Possible CSE 237D Project\)](#)

## Motivation

The code aims at automatic modulation classification for radio signals leveraging the RF capabilities offered by Xilinx ZCU111 RFSoc platform. The original code was created using python 3.6 and tensorflow 1.14 using the Vgg10 architecture. The network was trained using the radio modulation dataset created by O'Shea et. al available at [deepsig.io](https://deepsig.io).

The C code (trained neural network layers) was generated using a python script and implemented the quantized, low precision CNN components to mimic the precision of Verilog.

The top function `compute_network()` takes an input image signal of 1024 I/Q samples of the form `[I0,Q0,I1,Q1,..., I1023, Q1023]`. Output is a 4x1 vector of integers. The model is trained to detect 4 different modulation methods. We apply the softmax function to the output vector to choose between 1 of 4 modulation methods: OOK (0), 4ASK (1), 8ASK (2), BPSK (3). The softmax is the index where the maximum value occurs in an array.

We check the accuracy of the output against a validation set for different SNRs, and record the accuracy ( $100 \times \text{correct choices} / \text{incorrect choices}$ ) in output.txt.

## Testing With Python or HLS

### Code under test/

We improved upon the given code at [https://github.com/da-steve101/radio\\_modulation](https://github.com/da-steve101/radio_modulation). Under [https://github.com/da-steve101/radio\\_modulation/tree/master/c](https://github.com/da-steve101/radio_modulation/tree/master/c), their testing with python does not work with the deepsig.io dataset (they parsed it in their own way and did not link to the file) and is not documented.

### DEPENDENCIES:

- tensorflow
- python3
- debian machine for bash scripts

SETUP (for testing): In linux, download tar.gz from here: [deepsig.io](#) to code/

```
cp tw_vgg10_hardware.c tw_vgg10.c
bash get_txt_files.sh
export MODEL_DIR=../models/vgg_twn_rfsoc_50k_64_d128
rm -rf build
python3 setup.py build
export PYTHONPATH=build/lib.linux-x86_64-3.6
python3 test_vggpy.py
```

Or run:

```
cp tw_vgg10_hardware.c tw_vgg10.c; bash run_all.sh
```

For each snr in the validation set, accuracy is reported output.txt

For the python testbench, the output should be:

```
accr[16] = 25.714285714285715
```

```
...
```

```
accr[-14] =  
23.529411764705884  
  
accr[28] = 17.857142857142858
```

Where, for example, at SNR=-14 the accuracy is 23.529.

In Vitis HLS and in hardware emulation on Vitis V++ compiler, we feed a single 2048 input array. The output should *a/ways* display:

FAILED: (0,0) where **-186!**=210

FAILED: (0,1) where **199!**=144

FAILED: (0,2) where **34!**=-144

FAILED: (0,3) where **22!**=46

We used a legacy testbench from here

[https://github.com/da-steve101/radio\\_modulation/blob/master/c/tw\\_vgg10.c](https://github.com/da-steve101/radio_modulation/blob/master/c/tw_vgg10.c), and use wrong values to compare against the correct values (LHS, in bold above). In this way, we can make sure that the first 4 values are always the same. We also need to make sure that the next values are 0. Note that it would choose modulation classification #1 in this case. There are many output vectors that would provide a false negative, and this is one arbitrary way we chose to test our implementation.

# Initial Implementation with Vitis HLS

Note: Vitis HLS is found in \$(Vitis\_install\_path)/2019.2/bin/vitis\_hls. Its newer version of vivado HLS, and allows us to create .xo files after synthesis that can be used on the Vitis Unified Software Platform.

## Baseline (In Vitis HLS):

### **Code under baseline/**

The baseline folder contains the following files:

- src folder:
  - Top file: tw\_vgg10.c (and header file tw\_vgg10.h)
  - Helper files: all CNN layer (.c and .h) such as conv1.c, dense1.c
  - Testbench file: pyvgg.c
  - Script: script.tcl
- Report folder: contains the synthesis .rpt and .xml

Target Board: ZCU111 (part: xczu28dr-ffvg1517-2-e)

The generated C code from python implementation was not synthesizable. The use of pointer structures was throwing a lot of errors on synthesizing the code. We transformed the code to get a synthesizable baseline implementation in Vitis HLS.

Performance and resource utilization for our baseline implementation are given below.

The baseline implementation has a frequency of 132.1 Hz. As seen from resource utilization, the LUTs exceeds the resource utilization by almost 100%. This is because of the adders/subtractors in the CNN layers.

```

+ Timing:
  * Summary:
  +-----+-----+-----+
  | Clock | Target | Estimated | Uncertainty |
  +-----+-----+-----+
  | ap_clk | 10.00 ns | 8.750 ns | 1.25 ns |
  +-----+-----+-----+

+ Latency:
  * Summary:
  +-----+-----+-----+-----+-----+-----+
  | Latency (cycles) | Latency (absolute) | Interval | Pipeline |
  | min | max | min | max | min | max | Type |
  +-----+-----+-----+-----+-----+-----+
  | 103446 | 865110 | 1.034 ms | 8.651 ms | 103446 | 865110 | none |
  +-----+-----+-----+-----+-----+-----+

```

Fig- Performance of Baseline implementation

```

=====
== Utilization Estimates
=====
* Summary:
+-----+-----+-----+-----+-----+-----+
| Name | BRAM_18K | DSP48E | FF | LUT | URAM |
+-----+-----+-----+-----+-----+-----+
| DSP | - | - | - | - | - |
| Expression | - | - | 0 | 172 | - |
| FIFO | - | - | - | - | - |
| Instance | 422 | 474 | 125767 | 822673 | 0 |
| Memory | 423 | - | 0 | 0 | 0 |
| Multiplexer | - | - | - | 585 | - |
| Register | - | - | 151 | - | - |
+-----+-----+-----+-----+-----+-----+
| Total | 845 | 474 | 125918 | 823430 | 0 |
+-----+-----+-----+-----+-----+-----+
| Available | 2160 | 4272 | 850560 | 425280 | 80 |
+-----+-----+-----+-----+-----+-----+
| Utilization (%) | 39 | 11 | 14 | 193 | 0 |
+-----+-----+-----+-----+-----+-----+

```

Fig- Resource Utilization of Baseline implementation

# Optimizing the design in Vitis HLS

## Code under optimized/

The optimized folder contains the following files:

- src folder:
  - Top file: tw\_vgg10.c (and header file tw\_vgg10.h)
  - Helper files: all CNN layer (.c and .h) such as conv1.c, dense1.c
  - Testbench file: pyvgg.c
  - Script: script.tcl
- Report folder: contains the synthesis .rpt and .xml

Target Board: ZCU111 (part: xczu28dr-ffvg1517-2-e)

The following optimizations were used to improve the performance and resource utilization of our baseline design:

- Convolution functions and loops and array structure in file tw\_vgg10.c was simplified
- Dataflow implementation along with array partitioning (block factor=2) was performed for the design
- LUT Add/sub units were changed to DSP48E Add/Sub units in the CNN layers (conv2.c, conv3.c etc) to decrease our LUT usage
- The loops were pipelined and unrolled to increase throughput of our design

Performance and resource utilization for the optimized implementation are given below.

The optimized implementation has a frequency of 450 Hz which is almost 4x that baseline.

Our resource utilization of LUTs decreased significantly from baseline. However, it still exceeds the resource utilization by almost 83%. Even though resource allocation pragmas were used to shift some of the LUT adders/subtractors of the CNN layer files (conv2.c, conv3.c, dense1.c etc) to DSP48E adders/subtractors, the adders/subtractors of the vgg10 architecture could not fit the board using just the Vitis HLS pragmas.

```

=====
== Performance Estimates
=====
+ Timing:
  * Summary:
    +-----+-----+-----+-----+
    | Clock | Target | Estimated | Uncertainty |
    +-----+-----+-----+-----+
    | ap_clk | 10.00 ns | 8.750 ns | 1.25 ns |
    +-----+-----+-----+-----+

+ Latency:
  * Summary:
    +-----+-----+-----+-----+-----+-----+
    | Latency (cycles) | Latency (absolute) | Interval | Pipeline |
    | min | max | min | max | min | max | Type |
    +-----+-----+-----+-----+-----+-----+
    | 254338 | 254338 | 2.543 ms | 2.543 ms | 88102 | 88102 | dataflow |
    +-----+-----+-----+-----+-----+-----+

```

Fig- Performance of Optimized implementation

```

=====
== Utilization Estimates
=====
* Summary:
+-----+-----+-----+-----+-----+-----+
| Name | BRAM_18K | DSP48E | FF | LUT | URAM |
+-----+-----+-----+-----+-----+-----+
| DSP | - | - | - | - | - |
| Expression | - | - | 0 | 98 | - |
| FIFO | 0 | - | 646 | 3630 | - |
| Instance | 6 | 4240 | 212403 | 776350 | - |
| Memory | 258 | - | 0 | 0 | 0 |
| Multiplexer | - | - | - | 144 | - |
| Register | - | - | 16 | - | - |
+-----+-----+-----+-----+-----+-----+
| Total | 264 | 4240 | 213065 | 780222 | 0 |
+-----+-----+-----+-----+-----+-----+
| Available | 2160 | 4272 | 850560 | 425280 | 80 |
+-----+-----+-----+-----+-----+-----+
| Utilization (%) | 12 | 99 | 25 | 183 | 0 |
+-----+-----+-----+-----+-----+-----+

```

Fig- Resource utilization of Optimized implementation

# Running In Vitis

## **Code under vitis\_emulation/**

Now that we have optimized code, we run software emulation to check functionality, then hardware emulation in place of the physical chip. Hardware emulation defined by the Vitis toolset means that it emulates the target device. One of the main benefits of Vitis is that after hardware emulation, that is a high degree of confidence that it will run the same on chip due to the portability of the output file: xclbin. The target devices all run a linux cpu layer on top of an fpga, so portability isn't much of a problem.

The target device we use is **u200\_xdma\_201830\_2** in place of ZCU111, as the Vitis does not have support for ZCU111 as of yet. See [setup.pdf](#) for a walkthrough of how to install vitis and how to install the target device platform on ubuntu 18.04 if you are interested in running on your own machine, or want to expand on this project.



# Software and Hardware emulation outputs

The purpose of these outputs is to make sure functionality is correct. The outputs are the same as the vitis hls testbench output. There are also interesting logs showing that axis are being used in place of procedural code.

## Software Emulation (few minutes run time)

```
abc@abc-ThinkPad-P50:~/CSE237_FINAL_Ritika_Justin/vitis_cpp_run$ ./host compute_network.xilinx_u200_xdma_201830_2.xclbin
Found Platform
Platform Name: Xilinx
INFO: Reading compute_network.xilinx_u200_xdma_201830_2.xclbin
Loading: 'compute_network.xilinx_u200_xdma_201830_2.xclbin'
FAILED: (0,0)where -1861!=210
FAILED: (0,1)where 1991!=32
FAILED: (0,2)where 341!=144
FAILED: (0,3)where 22!=46
```

Hardware Emulation is most likely to work if software emulation works.

## Hardware Emulation (2hour+ run time)

```
abc@abc-ThinkPad-P50:~/CSE237_FINAL_Ritika_Justin/vitis_cpp_run$ ./host compute_network.xilinx_u200_xdma_201830_2.xclbin
Found Platform
Platform Name: Xilinx
INFO: Reading compute_network.xilinx_u200_xdma_201830_2.xclbin
Loading: 'compute_network.xilinx_u200_xdma_201830_2.xclbin'
INFO: [HW-EM 61] Hardware emulation runs simulation underneath. Using a large data set will result in long simulation times. It is recommended that a small dataset is used for
user verification. The flow uses approximate models for DDR memory and interconnect and hence the performance data generated is approximate.
Vitis Build version: 2.5.389
Build hash: 9a03798c11f86da2597b133db737c146a3ad84c
Build date: 2020-02-23 10:52:45
Git branch: 2019.2_PU2
V20: 307
UID: 1006
TTL: Mon 20 10:24:13 2020
```

```
HOST: abc-ThinkPad-P50
EXE: /home/abc/CSE237_FINAL_Ritika_Justin/vitis_cpp_run/host
[XRT] WARNING: unaligned host pointer '0x7ffe6b718ae0' detected, this leads to extra memcopy
[XRT] WARNING: unaligned host pointer '0x7ffe6b7189e0' detected, this leads to extra memcopy
INFO::[ Vitis-EM 22 ] [Time elapsed: 4 minute(s) 49 seconds, Emulation time: 0.287646 ms]
Data transfer between kernel(s) and global memory(s)
compute_network:m_axi_gmem-DDR[1] RD = 13.305 KB WR = 0.000 KB

INFO::[ Vitis-EM 22 ] [Time elapsed: 9 minute(s) 49 seconds, Emulation time: 0.579938 ms]
Data transfer between kernel(s) and global memory(s)
compute_network:m_axi_gmem-DDR[1] RD = 23.984 KB WR = 0.000 KB

INFO::[ Vitis-EM 22 ] [Time elapsed: 14 minute(s) 49 seconds, Emulation time: 0.935365 ms]
Data transfer between kernel(s) and global memory(s)
compute_network:m_axi_gmem-DDR[1] RD = 23.984 KB WR = 0.000 KB

INFO::[ Vitis-EM 22 ] [Time elapsed: 19 minute(s) 50 seconds, Emulation time: 1.31393 ms]
Data transfer between kernel(s) and global memory(s)
compute_network:m_axi_gmem-DDR[1] RD = 23.984 KB WR = 0.000 KB

INFO::[ Vitis-EM 22 ] [Time elapsed: 24 minute(s) 50 seconds, Emulation time: 1.66686 ms]
Data transfer between kernel(s) and global memory(s)
compute_network:m_axi_gmem-DDR[1] RD = 23.984 KB WR = 0.000 KB

INFO::[ Vitis-EM 22 ] [Time elapsed: 29 minute(s) 50 seconds, Emulation time: 2.05451 ms]
Data transfer between kernel(s) and global memory(s)
compute_network:m_axi_gmem-DDR[1] RD = 23.984 KB WR = 0.000 KB

FAILED: (0,0)where 3141!=210
FAILED: (0,1)where -23!=32
FAILED: (0,2)where -159!=--144
FAILED: (0,3)where 8!=46
INFO::[ Vitis-EM 22 ] [Time elapsed: 32 minute(s) 6 seconds, Emulation time: 2.24138 ms]
Data transfer between kernel(s) and global memory(s)
compute_network:m_axi_gmem-DDR[1] RD = 23.984 KB WR = 0.008 KB

INFO: [HW-EM 06-0] Waiting for the simulator process to exit
INFO: [HW-EM 06-1] All the simulator processes exited successfully
abc@abc-ThinkPad-P50:~/CSE237_FINAL_Ritika_Justin/vitis_cpp_run$
```

## Final Report files

After hardware emulation is run, final estimates are taken before being run on the board. The `compute_network.xilinx_u200_xdma_201830_2.xclbin` file in `vitis_emulation/` can be run on the sd card of the `xilinx_u200_xdma_201830_2` board.

The figures on the next few pages show the final estimates given by Vitis.

Figure 1

25	Timing Information (MHz)				
26	Compute Unit	Kernel Name	Module Name	Target Frequency	Estimated Frequency
27					
28	compute_network	compute_network	allocate_network	300.300293	430.663208
29	compute_network	compute_network	window_data_id_1	300.300293	411.015198
30	compute_network	compute_network	conv1	300.300293	421.22998
31	compute_network	compute_network	bn1_a_b	300.300293	420.698334
32	compute_network	compute_network	compute_conv_layer_1	300.300293	411.015198
33	compute_network	compute_network	window_data_id_2	300.300293	421.585175
34	compute_network	compute_network	conv2	300.300293	411.184235
35	compute_network	compute_network	bn2_a_b	300.300293	428.082184
36	compute_network	compute_network	compute_conv_layer_2_1	300.300293	411.184235
37	compute_network	compute_network	window_data_id_2_1	300.300293	419.463074
38	compute_network	compute_network	conv3	300.300293	411.184235
39	compute_network	compute_network	bn3_a_b	300.300293	430.477844
40	compute_network	compute_network	compute_conv_layer_2_2	300.300293	411.184235
41	compute_network	compute_network	window_data_id_2_2	300.300293	417.53653
42	compute_network	compute_network	conv4	300.300293	411.184235
43	compute_network	compute_network	bn4_a_b	300.300293	435.350464
44	compute_network	compute_network	compute_conv_layer_2_3	300.300293	411.184235
45	compute_network	compute_network	window_data_id_2_3	300.300293	415.627625
46	compute_network	compute_network	conv5	300.300293	411.184235
47	compute_network	compute_network	bn5_a_b	300.300293	435.350464
48	compute_network	compute_network	compute_conv_layer_2_4	300.300293	411.184235
49	compute_network	compute_network	window_data_id_2_4	300.300293	413.564941
50	compute_network	compute_network	conv6	300.300293	411.184235
51	compute_network	compute_network	bn6_a_b	300.300293	435.350464
52	compute_network	compute_network	compute_conv_layer_2_5	300.300293	411.184235
53	compute_network	compute_network	window_data_id_2_5	300.300293	411.692078
54	compute_network	compute_network	conv7	300.300293	411.184235
55	compute_network	compute_network	bn7_a_b	300.300293	432.900452
56	compute_network	compute_network	compute_conv_layer_2_6	300.300293	411.184235
57	compute_network	compute_network	dense1	300.300293	411.184235
58	compute_network	compute_network	bn1d1_a_b	300.300293	428.265503
59	compute_network	compute_network	dense2	300.300293	411.184235
60	compute_network	compute_network	bn2d2_a_b	300.300293	428.082184
61	compute_network	compute_network	compute_network	300.300293	411.015198
62					

Figure 2

[illegible]



Figure 3

Area Information	Kernel Name	Module Name	FF	LUT	DSP	BRAM	URAM
compute_network	compute_network	allocate_network	715	1159	3	1	0
compute_network	compute_network	window_data_id_1	1189	850	8	0	0
compute_network	compute_network	conv1	756	3712	0	0	0
compute_network	compute_network	bn1_a_b	2178	4280	59	0	0
compute_network	compute_network	compute_conv_layer_1	4510	9494	71	0	0
compute_network	compute_network	window_data_id_2	998	744	8	0	0
compute_network	compute_network	conv2	19111	52167	0	0	0
compute_network	compute_network	bn2_a_b	2042	4715	45	0	0
compute_network	compute_network	compute_conv_layer_2_1	22452	58231	56	0	0
compute_network	compute_network	window_data_id_2_1	997	744	8	0	0
compute_network	compute_network	conv3	18834	52029	0	0	0
compute_network	compute_network	bn3_a_b	2010	4702	45	0	0
compute_network	compute_network	compute_conv_layer_2_2	22128	58203	54	0	0
compute_network	compute_network	window_data_id_2_2	996	744	8	0	0
compute_network	compute_network	conv4	18687	51453	0	0	0
compute_network	compute_network	bn4_a_b	1954	4437	51	0	0
compute_network	compute_network	compute_conv_layer_2_3	21911	57329	60	0	0
compute_network	compute_network	window_data_id_2_3	1035	832	8	0	0
compute_network	compute_network	conv5	18422	51479	0	0	0
compute_network	compute_network	bn5_a_b	1893	4377	51	0	0
compute_network	compute_network	compute_conv_layer_2_4	21613	57376	60	0	0
compute_network	compute_network	window_data_id_2_4	1034	832	8	0	0
compute_network	compute_network	conv6	17481	50570	0	0	0
compute_network	compute_network	bn6_a_b	1813	4024	57	0	0
compute_network	compute_network	compute_conv_layer_2_5	20579	56108	66	0	0
compute_network	compute_network	window_data_id_2_5	1033	832	8	0	0
compute_network	compute_network	conv7	17559	50123	0	0	0
compute_network	compute_network	bn7_a_b	1772	4674	42	0	0
compute_network	compute_network	compute_conv_layer_2_6	20619	56305	51	0	0
compute_network	compute_network	densel	109736	297783	0	0	0
compute_network	compute_network	bnd1_a_b	1722	9002	50	0	0
compute_network	compute_network	dense2	25568	84482	0	0	0
compute_network	compute_network	bnd2_a_b	1896	8254	74	0	0
compute_network	compute_network	compute_network	274558	756918	546	551	0

Figure 4

The screenshot displays the Vivado IDE interface during a synthesis process. The left-hand pane shows a project hierarchy under 'BUILD'. The main window is divided into two sections: 'Synthesis Report for \'compute\_network\'' and 'Performance Estimates'.

### Synthesis Report for 'compute\_network'

**Vivado HLS Report for 'compute\_network'**

- Date: Fri Mar 20 20:18:43 2020
- Version: 2019.2 (Build 2704478 on Wed Nov 06 22:10:23 MST 2019)
- Project: compute\_network
- Solution: solution
- Product family: virtexplus
- Target device: xcu200-fsgd2104-2-e

### Performance Estimates

#### Timing

**Summary**

Clock	Target	Estimated	Uncertainty
ap_clk3.33 ns	2.433 ns	0.90 ns	

#### Latency

**Summary**

**Detail**

**Instance**

**Loop**

### Utilization Estimates

**Summary**

Name	BRAM_18K	DSP48E	FF	LUT	URAM
DSP	-	1	-	-	-
Expression	-	-	0	168	-
RIFO	-	-	-	-	-
Instance	3	545	274200	754780	-
Memory	548	-	0	0	0
Multiplexer	-	-	-	1970	-
Register	-	-	358	-	-
Total	551	546	274558	756918	0
Available	4320	684023644801	182240	960	
Available SLR	1440	2280	788160	394080	320
Utilization (%)	12	7	11	64	0
Utilization SLR (%)	38	23	34	192	0

**Detail**

**Instance**

If the run is successful, you should get two classes of report files:

- Get report files from hardware emulation (this runs vivado hls or vitis hls).
  - Run vitis\_analyzer in the <Vitis\_install\_location>/2019.2/bin and open the link summaries to get the Estimated Timing Information (Figure 1), Latency Information (Figure 2) and Area Information (Figure 3).
  - Open the Compile Summary to get the estimated latency and area on the board (figure 4).

**Figure 4 shows that the design fits for our target device - u200\_xdma\_201830\_2.**

Although, it estimates that the SLR portion of the board is overutilized. This is just an estimate, and we can't change this without running board specific optimizations (which is counterproductive because we want to eventually migrate to ZCU111) [1].

[1]:

[https://www.xilinx.com/html\\_docs/xilinx2019\\_1/sdaccel\\_doc/xqd1554752221979.html#p hm1523550940155](https://www.xilinx.com/html_docs/xilinx2019_1/sdaccel_doc/xqd1554752221979.html#p hm1523550940155).

# Tradeoffs between Vitis and Vivado

Figure A

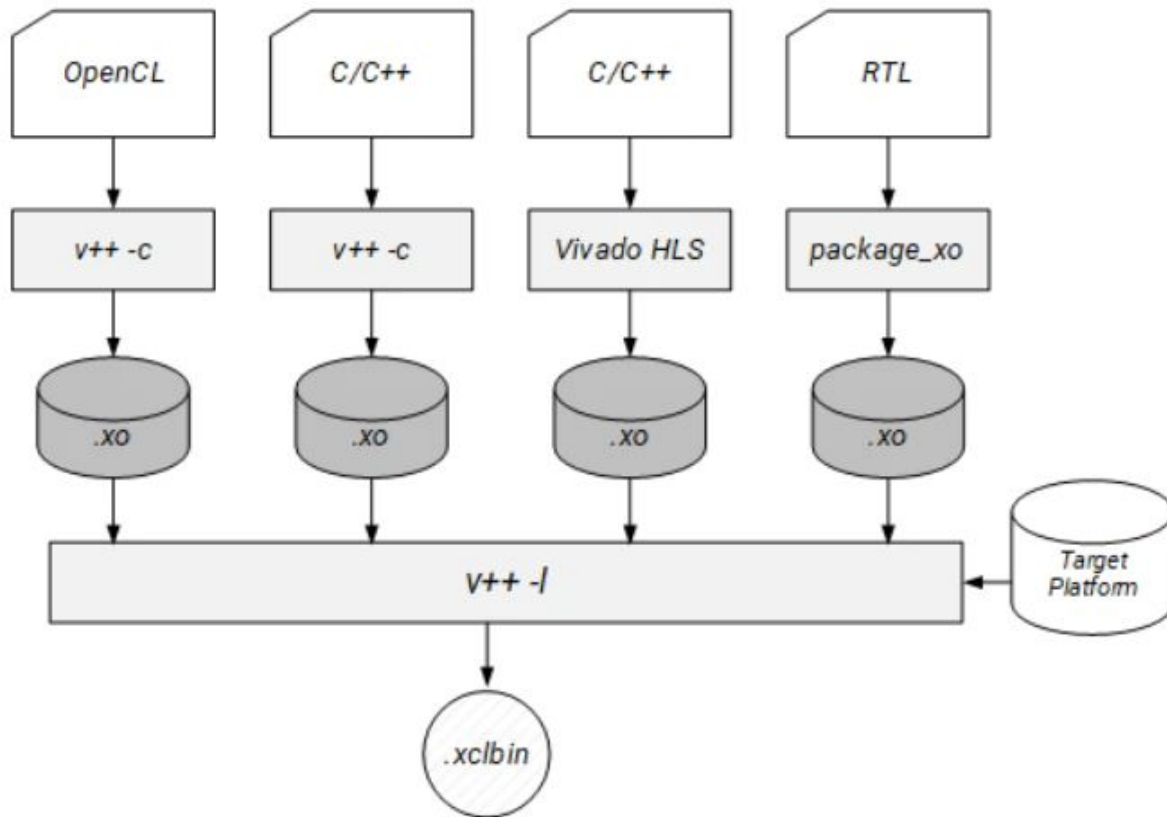


Figure A is from <https://www.hackster.io/news/microzed-chronicles-vitis-hls-85a9b1d272e6>

## Pros of the Vitis Unified Software Platform

- Uses host file and a kernel file
- See diagram. Kernel is “plug-and-play” i.e. it can be linked with many programming languages for the host file.
- Allows for multiple programming languages. C for the kernel. The host can be cpp, c or python.
- See the flow below, it allows us to test in software with axis, before running on hardware. Our design was medium sized, and it took around 2 hours to run the Hardware steps below. If our design was large, this would be an even bigger benefit.
- Platforms that are supported are easy to download (the dpkg files are given for ubuntu users).

## Cons of the Vitis Unified Software Platform

- Very little documentation
- Still very buggy. For instance, although it supports C, there is a bug that using a c file extension will not allow vivado hls to export the .xo file. This can only be fixed by changing c files to cpp.
- Only really supports linux. There is a gui for windows, but it is documented even less than the ubuntu version. Had to use ubuntu 18.04 to run the tutorials.
- Accelerated Platforms (i.e. platforms that get the upsides mentioned above) are limited for now.

## Future Work (Possible CSE 237D Project)

- Run this code on the ZCU111 board using the PYNQ interface, probably have to do vivado as ZCU111 is not supported officially by vitis yet.
- Modify the axi interface to stream 2048 inputs one at a time
- Optimize in vitis, the possibilities are much more than we were able to go though. Memory buffers are the most obvious route to take.
- Optimize in hls, as we learned in class there we can explore further the pareto curve. We made many sacrifices for the area (this machine learning model was 10 layers of 1000-input look up tables) but this was not enough for ZCU111.
- Suggested by Alireza - Try a different architecture, such as imagenet, that is larger but gives better accuracy.